RTX Project Report

Xiang, Dian
20431601
dxiang@uwaterloo.ca

Justin McGirr
20413625
jmcgirr@uwaterloo.ca

Adrian Cheung
20421743
a32cheun@uwaterloo.ca

Aaron Morais
20413440
aemorais@uwaterloo.ca

April 6, 2014

Abstract

Contents

List of Algorithms

List of Figures

Part I

Introduction

Part II

Awesome RTX

Chapter 1

Global Variables

1.1   Description

There are various amounts of variables used throughout the RTX. This section describes the global variables that are used to accomplish tasks in the RTX. Other global variables such as the global process IDs can be seen in the appendix TODO. There are three major sections that required global variables and data structures: memory management in the heap, scheduler, and user processes.

1.2   Heap Data Structures

The RTX provides functionality to request and release memory from the heap, which is shared and stored in the RAM of the board. There is one main data structure that stores each memory block.

```c
typedef struct HeapBlockHeader {
    int source_pid;
    int dest_pid;
    unsigned int send_time;
    struct HeapBlock* p_next;
    struct HeapBlock* p_next_usr;
} HeapBlockHeader;

#define HEAP_BLOCK_SIZE 128

typedef struct HeapBlock {
    HeapBlockHeader header;
```

```
13    byte              data[HEAP_BLOCK_SIZE];
} HeapBlock;
```

The HeapBlock structure stores a header and the content of the block. When a memory is requested, users are given the data with a adjustable size of HEAP_BLOCK_SIZE and does not have knowledge of the header. Helper functions in the RTX turn the user block back into kernel block by adjusting the pointer of the block. Each header contains information about message passing (if the block is used for this purpose) and a pointer to the next block (for kernel memory management purposes). This global data structure is used in the following sections:

1. **Memory Management** - this section uses this data structure to give and receive back heap memory blocks

2. **Process Message Passing** - this section uses this data structure pass messages between processes. Message envelope are implemented as HeapBlocks.

3. **Timer I-Process** - The process uses message passing to send delayed messages.

4. **UART I-Process** - The process uses message passing to process input and output characters.

5. **CRT I-Process** - The process receives messages and passes message for terminal output.

6. **KCD Process** - The KCD uses message passing to CRT for display and user processes for processing.

1.3   Process Scheduler Data Structures

The RTX has a fixed-priority based scheduler that acts as a uniprocessor system. Context switching is required and the following data structures are used for this purpose.

1.3.1 Ready and Blocked Priority Queues

On initialization, each in memory process is given a process control block
(PCB). The PCB contains data on the process such as its process ID, stack
pointer, process state, and priority that the kernel will use for scheduling.
The blocked and ready priority queues of PCBs keep track of which processes
are blocked and ready for execution. Each process is in one of the 5 states at
all time listed in ProcessState. More details about the states can be found
in section TODO. Processes that are in state PROCES_STATE_READY are
on the ready queue. Processes in state PROCESS_STATE_BLOCKED are
in the blocked queue. There is also an implicit queue for processes blocked
on message but we won't go into that in this section.

```
1 typedef enum {
    PROCESS_STATE_NEW                 = 0,
3   PROCESS_STATE_READY               = 1,
    PROCESS_STATE_RUNNING             = 2,
5   PROCESS_STATE_BLOCKED             = 3,
    PROCESS_STATE_BLOCKED_ON_MESSAGE  = 4,
7 } ProcessState;

9 typedef struct PCB {
    // Stack pointer
11  U32*            sp;

13  ProcessID       pid;
    ProcessState    state;
15  ProcessPriority priority;

17  struct PCB*     p_next;
    // Incoming messages, waiting to be processed.
19  HeapBlock*      message_queue;
  } PCB;
21
  PCB* g_ready_process_priority_queue[PROCESS_PRIORITY_NUM] = {
      NULL};
23 PCB* g_blocked_process_priority_queue[PROCESS_PRIORITY_NUM];
```

PCBs and the priority queues are used in the following sections:

1. **Memory Management** - When a process requests memory without
   available memory blocks, the memory management unit must add the
   process to a blocked queue. When a memory block is released, the pro-

cess with the highest priority in the blocked queue (if any) is moved into the ready queue. Thus, the memory management unit must have access to the PCBs, and blocked and ready queues in order to accomplish these task.

2. **Process Management** - The process management unit is the one that schedules processes and keep track of the process priorities. Ready processes are taken from the ready queue to be ran if the current process is blocked or finished its execution. The process management unit also takes care of updating the two priority queues when the priority of a process has been changed.

3. **UART Process** - The UART Process needs access to the PCB. The PCBs also contain a mailbox for messages. The UART displays any incoming messages to the terminal display. Thus, it needs to know the PCB structure in order to gain access to the mailbox.

1.3.2   Priorities

The RTX is based on a fixed-priority scheduler with 5 user process priorities and 2 system process priorities given below. Users are given priorities PROCESS_PRIORITY_HIGH to PROCESS_PRIORITY_LOWEST. PROCESS_PRIORITY_NULL_PROCESS is given to the null process and PROCESS_PRIORITY_SYSTEM_PROCESS are given to critical processes such as the KCD, CRT, timer I-process, and UART I-process.

```
1  typedef enum {
      PROCESS_PRIORITY_INVALID          = 0 ,
3     PROCESS_PRIORITY_SYSTEM_PROCESS   = 1 ,
      PROCESS_PRIORITY_HIGH             = 2 ,
5     PROCESS_PRIORITY_MEDIUM           = 3 ,
      PROCESS_PRIORITY_LOW              = 4 ,
7     PROCESS_PRIORITY_LOWEST           = 5 ,
      PROCESS_PRIORITY_NULL_PROCESS     = 6 ,
9     PROCESS_PRIORITY_UNSCHEDULABLE    = 7 ,

11    PROCESS_PRIORITY_NUM              = 8
   }  ProcessPriority ;
13
   typedef enum {
15    USER_PROCESS_PRIORITY_HIGH                = 0 ,
      USER_PROCESS_PRIORITY_MEDIUM              = 1 ,
```

```
17    USER_PROCESS_PRIORITY_LOW              = 2,
      USER_PROCESS_PRIORITY_LOWEST          = 3,
19
      USER_PROCESS_PRIORITY_NUM             = 4,
21  } UserProcessPriority;
```

This priority structure is used by the process management unit to schedule and block processes based on their priorities. The RTX provides users with UserProcessPriority while keeping an internal structure of ProcessPriority.

1.3.3   Message Passing

Message passing is a way for interprocess communication provided by the RTX. Messages are created in the form of a msgbuf structure, which includes the type and content. Some processes such as the CRT and KCD will require a certain type of message to be sent before the message can be processed correctly.

```
1  typedef enum {
     MESSAGE_TYPE_KCD_KEYPRESS_EVENT          = 0,
3    MESSAGE_TYPE_KCD_COMMAND_REGISTRATION = 1,
     MESSAGE_TYPE_CRT_DISPLAY_REQUEST         = 2,
5    MESSAGE_TYPE_WALL_CLOCK                   = 3,
     MESSAGE_TYPE_COUNT_REPORT                = 4,
7    MESSAGE_TYPE_WAKEUP_10                    = 5,

9    MESSAGE_TYPE_NUM                          = 6,
   } MessageType;
11
   struct msgbuf {
13   MessageType mtype:
     char mtext[HEAP_BLOCK_SIZE - sizeof(MessageType)];
15 };
```

The message passing data structure is evident in the KCD, wall clock process, and all I-processes. It can also be used in all user processes.

Chapter 2

Kernel API

2.1 Description

This section describes the kernel API that is available to users in the RTX. It will only go into details of how to use each API function call and states of different scenarios. Details of the implementation can be found in section TODO or from the raw code.

2.2 Memory Management

2.2.1 Description

The RTX provides the utility of simple memory management of the heap. The main memory on the board (TODO: name the board) is divided into sections of the RTX Image, PCB data, the heap, and process stacks seen in figure blah. *** TODO insert figure blah here ***

Board LPC17xx (TODO insert the real name) does not have the hardware to support virtual memory. Thus, a fixed memory management scheme is used with variable block sizes. The OS kernel is alway loaded into main memory. When the OS boots, a user stated number of PCBs and process stack are allocated. The OS also allocates memory for system processes such as the Keyboard Command Decoder Process (KCD) and the CRT Display Process. After allocation of PCBs and process stacks, the remaining memory is used for the heap, which is shared between all processes. This section will be focused on the memory management of the heap.

The heap is further divided into a variable number of blocks. Each block contains a header (HeapBlockHeader) and 128 bytes of data. Depending on the number of user and system processes, the number of available heap memory will vary. The size of the data can vary from 128 bytes but must be set at compile time. The OS supports two kernel calls that gives access to these memory blocks.

2.2.2   Requesting Memory

```
void *request_memory_block();
```

The first functionality supported by the OS is the ability to request memory. This call gives back a pointer to a memory block in the heap. The size of the memory block is defined by the constant HEAP_BLOCK_SIZE, which defaults to 128 bytes. These blocks are used for storing local variables or as envelopes for interprocess communication (described in section TODO). The user must cast the memory block to the proper type for his or her own use. An example of the usage is provided below, which stores the numbers 0-31 in an array of size 32.

```
1  void user_process() {
       int size = 32;
3
       int* array = (int*)request_memory_block();
5
       for( int i = 0; i < size; i++ ) {
7          array[i] = i;
       }
9
       // ...
11
       release_memory_block( (void*)array );
13 }
```

All processes will share the same heap memory pool. Thus, this primitive will block the process if the OS does not have anymore memory blocks to give out at the time, thus effecting the execution of the process. In that case, it will only be unblocked there is a new memory block available and if it has the highest priority on the list of processes waiting for a memory block. When using the memory block, the user must be aware of writing past the

heap block size. The OS does not check for segmentation faulting. Thus, undefined behavior may occur. Also, the user must remember to release the memory block or memory leakage will occur.

Requesting memory is reliant on calls to the heap to find an available memory block to allocate. The heap keeps a list of the status of each memory block to determine whether or not a block of memory can be requested. The following is the implementation used by the heap:

function FINDFREEBLOCK
    for $i = 0$ to $TotalNumberOfBlocks$ **do**
        if FreeSpaceBitmap[i] is a free block **then**
            return i
        end if
    end for
    return not found
end function

The lookup time of the FreeSpaceBitmap is O(1) as it is a simple array lookup and a boolean evaluation. Finding a free block has its dependencies on the total number of blocks. The best case for requesting memory is O(1) when the first block in the heap is free. Since this is a linear search, on average the complexity is O(n) where n is the total number of blocks.

2.2.3   Releasing Memory

```
1 int release_memory_block(void *memory_block);
```

The second functionality supported by the OS is the ability to release memory. This is a non-blocking call that returns the memory block back to the OS. This should be called when a message is received and not passed on or when the process is done using the memory block. If any process is blocked on memory, it will be unblocked and put to the ready queue. If the current process has lower priority than the unblocked process, then the current process will be preempted and the higher priority process will be executed instead. An example of this can be seen in figure CODE1 TODO on line 13.

In order to release the memory, release_memory_block will need to rely on the implementation of the heap. As stated in 2.2.2, the heap keeps a list

15

of all available memory blocks. In order to release a memory block, the heap
will set that block in the list as a free block. Releasing the block thus has
O(1) complexity as it only relies on an array lookup. The implementation is
as follows:

    function RELEASEMEMORY(MemoryBlock)

        FreeSpaceBitmap[position of MemoryBlock] ← true

    end function

2.3   Processor Management

```
int release_processor();
```

The OS manages processes as though it is on a uniprocessor. A priority
based scheme with context switching is used for scheduling processes. A
process can voluntarily release the processor to the OS at any time during
its execution. If there are errors in the call, it will return an error without
releasing the processor to the OS. If there are no errors and the invoking
process is ready to execute, it is put to the end of the ready queue of its
priority. If there are no other processes of equal or higher priority, the process
will be chosen to execute again. However, if there is, another process may be
chosen for execution. Below is an example which prints an increasing number
and releases the processor at every turn.

```
void user_process() {
    static num = 0;
    while (1) {
        print(num++);
        release_processor();
    }
}
```

2.4   Interprocess Communication

Apart from heap memory, processes do not share information. Thus, com-
munication between processes is done through message-based interprocess
communication (IPC). Details of the internal layout of process mailboxes

can be found in section TODO. The RTX gives three primitives to carry out this task, one for sending, one for delayed sending, and one for receiving.

2.4.1   Send Messages

```
1  int send_message(int process_id, void *message_envelope);
```

A process can compose a message envelope to be sent to another process. Memory for envelope message must be requested from the RTX using the request_memory() routine. The envelope consists of a type (mtype) and the message data (mtext) which must be filled in as seen below. The predefined message types are used for the KCD, CRT, and Wall Clock process, which are built into the RTX (see section TODO for more information). The message data has a predefined size which is a MessageType smaller than the HEAP_BLOCK_SIZE. Thus, any message that are longer will exhibit undefined behavior. The process ID of the receiving process must also be known ahead of time in order to use send_message.

```
1  typedef enum {
       MESSAGE_TYPE_KCD_KEYPRESS_EVENT           = 0,
3      MESSAGE_TYPE_KCD_COMMAND_REGISTRATION  = 1,
       MESSAGE_TYPE_CRT_DISPLAY_REQUEST          = 2,
5      MESSAGE_TYPE_WALL_CLOCK                    = 3,
       MESSAGE_TYPE_USER_DEFINED                 = 4,
7
       MESSAGE_TYPE_NUM                          = 5,
9  } MessageType;

11 struct msgbuf {
       MessageType mtype;
13     char mtext[HEAP_BLOCK_SIZE - sizeof(MessageType)];
   };
```

The primitive returns a status which validates the, message, receiver process ID, and ready queue. User processes is allowed to send a message to any user processes or system processes such as the KCD. If the receiving process is currently blocked on receiving message, this call will add the message to the process' message box and put it back on the message queue. The current process continues to execute unless the receiving process has a higher

17

priority. In that case, the current process will be preempted and put to the back of the ready queue. Thus, this primitive may effect the execution of the process.

In order to keep track of messages that have been sent, the messages are stored in a queue, specifically a linked list. Messages are added to the end of the linked list resulting in an O(n) insertion complexity. It is possible to have this run in O(1) time by having the inserting take place at the front of the linked list, however this is a trade-off for receiving messages. In the current implementation, received messages will be taken from the front of the linked list having O(1) complexity. If the change was to be made to optimize sending messages, then this would slow down receiving messages to O(n).

2.4.2   Receive Messages

```
void *receive_message(int *sender_id);
```

A process can use this primitive to receive messages from other processes. Unless the sender_id is NULL, the sender of the message will be written to the sender_id parameter. The current process will check its mailbox for any incoming messages. If there are no messages in its mailbox, the routine will block the process and select another process for execution. Execution of the process will occur again if a message is sent from another process and this process has the highest priority in the ready queue. The message_envelope are heap memory underneath. Thus, it is required that unless the process passes the message onto another process, the final receiving process must release the message_envelope block after its usage. An example of send_message and receive_message is shown below.

```
1  void process_send() {
       while(1) {
3          // Request one block for process_receive to use
           struct msgbuf* message_envelope
5              = (struct msgbuf*)request_memory_block();
           message_envelope->mtype = MESSAGE_TYPE_USER_DEFINED_1;
7          strcpy(message_envelope->mtext, "process message");
           send_message(PROCESS_RECEIVE_ID, message_envelope);
9
           // Send another block for process_receive to send to CRT
11         struct msgbuf* message_envelope
               = (struct msgbuf*)request_memory_block();
```

18

```
13          message_envelope->mtype = MESSAGE_TYPE_USER_DEFINED_2;
            strcpy(message_envelope->mtext, "print this");
15          send_message(PROCESS_RECEIVE_ID, message_envelope);
        }
17 }

19 void process_receive() {
       while(1) {
21          struct msgbuf* message
                = (struct msgbuf*)receive_message(NULL);
23          if( message->mtype = MESSAGE_TYPE_USER_DEFINED_2 ) {
                // Send to CRT for printing; do not release memory
25              message->mtype = MESSAGE_TYPE_CRT_DISPLAY_REQUEST;
                send_message(PROCESS_ID_CRT, message);
27          } else {
                // Do something with the message then release it
29              release_memory_block( (void*)message );
            }
31      }
   }
```

2.4.3   Delayed Messages

```
int delayed_send(int process_id, void *message_envelope, int
    delay);
```

This primitive is very similar to the primitive of send_message in section 2.4.1 with the addition of a delay. Instead of sending the message immediately, the message will be sent delay number of milliseconds. The message_envelope is constructed in the same way. The process_id will be the receiving process. The execution of the process may be preempted after a the delay period if the receiving process was blocked on receive_message and has a higher priority.

2.5   Process Priority

The RTX schedules processes based on a fixed priority scheme. Thus, the RTX provides the utility to change priorities and to get the process priority of any user processes.

2.5.1   Set Process Priority

```
int set_process_priority(int process_id, int priority);
```

This primitive allows user processes to set the priorities of other user processes. User processes are not allowed to set the priorities of any system processes. System processes, however, are allowed to set the priorities of other system or i-processes. The primitive validates if the current process is allowed to set the priority of the process with process_id and whether it's a valid process_id. It also checks the validity of the priority. An error status of RTX_ERR is given back if the validation does not pass. If properly set, a status of RTX_OK is returned. The caller of this primitive is not blocked but can be preempted if the priority of the set process is higher. Otherwise, the process continues to execute. An example usage is shown in section 2.5.2

2.5.2   Get Process Priority

```
int get_process_priority(int process_id);
```

Given a process_id, this primitive gives back the priority of any processes including any system or i-processes to a user process. A invalid process_id will result in a RTX_ERR status. Otherwise, a RTX_OK is returned to the caller. An example usage of both get and set process priority is shown below.

```
// Assume this process has a process ID of 3
// and there is another user process with ID of 1
void user_process() {
    int process_priority_3 = get_process_priority(3);
    if( process_priority_3 == USER_PROCESS_PRIORITY_MEDIUM ) {
        // Make process 1 have a higher priority
        // This call will preempt
        int status = set_process_priority(
                          1, USER_PROCESS_PRIORITY_HIGH );
        if( status == RTX_ERR ) {
            release_processor();
        }
    }
}
```

Chapter 3

Interrupts and Their Handler and Processes

Chapter 4

System and User Processes

4.1 System Processes

4.1.1 Description

The RTX consists of system processes that sit on top of the kernel. These system processes usually have higher priority than user processes because they carry out tasks for the OS that the kernel cannot do otherwise. System processes in this RTX includes the null process, and I/O processes.

4.1.2 Null Process

The design of the RTX requires that a process should always be executing. Thus, if there are no user or other system processes, the null process executes. The only job of the null process is to keep the processor running. It has the lowest priority of any process using the RTX. If any process is added into the ready queue during the execution of this process, a preemption will occur and the new process will take control of the processor. The final implementation follows the following form.

```
function NULL_PROCESS
    while true do
    end while
end function
```

4.1.3 CRT Display Process

The RTX supports output to terminal display. CRT is a system process that the RTX uses to accomplish the task of display. The CRT process depends on the UART I-Process (details of the UART can be found in section TODO). The CRT is implemented using the following algorithm.

function CRT_PROCESS
    while true **do**
        message = RECEIVE_MESSAGE(NULL)
        if message is not a CRT display request **then**
            continue
        end if
        SEND_MESSAGE(PROCESS_ID_UART, message)
        Set interrupt bits
    end while
end function

The CRT display process accepts message requests from any process of type MESSAGE_TYPE_CRT_DISPLAY_REQUEST. This tells the CRT to send the message for display to the UART I-process, which interacts with the hardware registers to output to terminal. Since the UART process is an interrupt process, the interrupt bit must be set in order for the process to run. Thus, after sending the message to UART's mailbox, the CRT process enables the interrupt bits for the I-process to trigger. The UART I-process passes the message character by character to the hardware to display to the terminal.

4.1.4 Keyboard Command Decoder Process

The RTX provides functionality to forward user input commands to a particular process. The keyboard command decoder process (KCD) handles two tasks: giving input to the CRT for terminal display and forwarding certain inputs to different processes for further processing. For the first task, all keyboard inputs are filtered by the KCD. The UART I-process is called when a key is first pressed. The UART I-process (explained further in section TODO) then passes a message to the KCD telling of type key press. In the first case, the KCD will take this message and pass it on to the CRT for terminal display.

In the latter case, processes can register a command with the KCD. Each

command registration consists of a single capital letter of the alphabet. If a command is taken, the KCD will not register the command and exit with an error. After a valid command has been registered, user input with the pattern "%" followed by a upper case character will be parsed by the KCD. If user types in "%" followed by a registered character, the KCD will buffer the input until the user presses the return key. In that scenario, the KCD will pass it onto the registered process. The following algorithm of the KCD explains this in more detail.

function PROCESS_INPUT(message)
    if There a "%" character in the buffer **then**
        if message.Content is a command character **then**
            add to buffer
        else
            clear buffer
        end if
    else if message.Content is "%" **then**
        add to buffer
    else if message.Content is enter and buffer is not empty **then**
        send buffer to registered command
    end if
    Send message to CRT to display
end function

function PROCESS_REGISTER(message)
    if message.Command has not been taken **then**
        register message.Command
    else
        error
    end if
end function

function KCD_PROCESS
    initialization
    while true **do**
        message = RECEIVE_MESSAGE(NULL)
        if message type is key press input **then**
            PROCESS_INPUT(message)
        else if message type is command registration **then**

PROCESS_REGISTER(message)
        end if
      end while
   end function

An example of this usage is the wall clock process (described in more detail in section 4.2.1). The wall clock process registers the character 'W' as a command. The registration will be done by passing a message to the KCD of command registration type. This message is received by the KCD. If the 'W' command has not been registered, it becomes registered by the wall clock process. Meanwhile, the user types in random characters. Any character that is not '%' is passed by the KCD to the CRT for display. Once a '%' is received, the KCD checks the next character to see if it's a registered command. If its not a registered command, the KCD clears the buffer and starts at the beginning. If it is a registered command, the KCD keeps a buffer of all the input characters. This buffer fills with data and is sent to the wall clock process for processing after the user presses enter. The buffer is only 20 characters long. Thus, any data for a process after 20 characters will not be sent to the registered process on enter.

4.2   User Processes

4.2.1   Wall Clock Process

4.2.2   'funProcess'

4.2.3   'schizophrenicProcess'

4.2.4   'fibProcess'

4.2.5   'memoryMuncherProcess'

4.2.6   'releaseProcess'

Chapter 5

Initialization

5.1   Description

5.2   Memory Layout

5.3   Process Mailbox

Chapter 6

Major Design Changes

Part III

Testing and Analysis

Chapter 7

Testing

7.1   Description

7.2   Theoretical Analysis

7.3   Measurements

Chapter 8

Timing Analysis

8.1 Requesting and Releasing Memory

All data used and processed data can be found in the Appendix in A

8.1.1 Single Request Single Release

In the first analysis, the request_memory and release_memory methods were
timed. Single request, single release means that a single memory block was
requested and the time taken was recorded. Similarly for releasing a memory
block. By doing so we are able to keep the heap in the same state, eliminating
the effects of cumulative memory requests (which will be tested in section
8.1.2). The method to time this is shown in the following algorithm. Data
was collected across 5 machines, with 1000 calls to each function.

   function SINGLE_REQUEST_SINGLE_RELEASE
      for $i = 0$ to 1000 **do**
         RequestStartTime[i] ← Curent Time
         memoryBlock ← REQUEST_MEMORY_BLOCK
         RequestEndTime[i] ← Curent Time
         ReleaseStartTime[i] ← Curent Time
         RELEASE_MEMORY_BLOCK(memoryBlock)
         ReleaseEndTime[i] ← Curent Time
      end for
   end function

Based on the analysis, the mean run time for release memory is 405.44ns. The standard error (marked in the green mean diamonds on the figure) is very small, giving a 95% confidence interval ranging from [404.99, 405.84] to [405.1, 405.95]. The F ratio is incredibly small (0.0474), meaning that there is little to no variance among the different machines and as such it can be concluded that the hardware does not greatly affect the performance of the method.
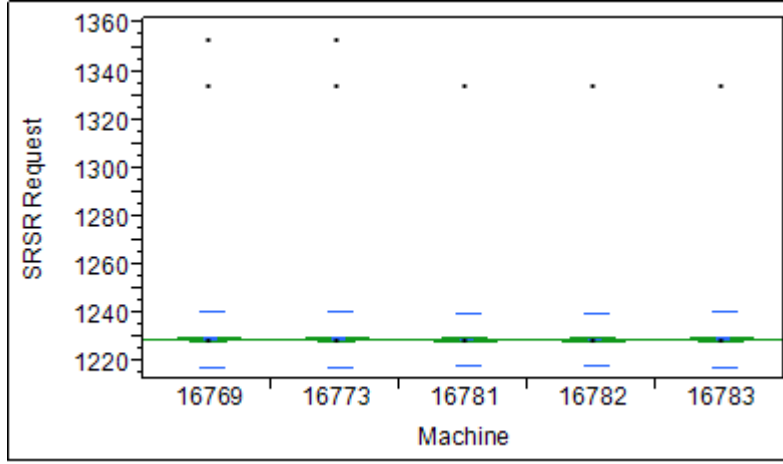
Figure 8.1: Request Memory over 5 machines

The results of releasing memory are very similar to that of requesting memory. The mean time taken to release memory is 1229.237ns. The variations among different machines were very small, giving an F Ratio of 0.0328. This shows that the performance is consistent across different hardware. The performance was as well consistent within itself, the fastest machine had a 95% confidence interval of [1228.5, 1229.9], and for the slowest [1228.6, 1230.0]. Releasing memory takes on average a little under three times as long as requesting memory

8.1.2 Cumulative Memory Requests

The cumulative memory requests condition had 50 memory blocks requested and then subsequently released. This differs from experiment in 8.1.1 where the requested memory blocks do not accumulate. The purpose of this experiment is to see the effect of depleting resources on run time. Once again,

31

Figure 8.2: Release Memory over 5 machines

trials were taken over 5 machines, however, each machine showed the same performance. To acheive this, the following two method calls were made.

function CUMULATIVE_REQUEST
 for $i = 0$ to 50 **do**
  RequestStartTime[i] ← Curent Time
  memoryBlock ← REQUEST_MEMORY_BLOCK
  RequestEndTime[i] ← Curent Time
 end for
end function
function CUMULATIVE_RELEASE
 for $i = 0$ to 50 **do**
  ReleaseStartTime[i] ← Curent Time
  memoryBlock ← RELEASE_MEMORY_BLOCK
  ReleaseEndTime[i] ← Curent Time
 end for
end function

Consistent with the Single Request, Single Release experiment, the first memory block requested took 1229ns. A linear trend line is overlaid on top showing an R2 value of 1 (a perfect trend). The trend shows that for each additional memory block requested, the run time increase by 17ns. This linear trend is in accordance with the implementation of the request memory method which has O(n) complexity described in 2.2.2.

32

Figure 8.3: Release Memory over 5 machines

8.2 Sending and Receiving Messages

8.2.1 Single Send, Single Receive

In this experiment, the send_message and receive_message methods were timed. A single message was sent and the time taken was recorded. Similarly, the time taken to receive that message was timed, specifically the time for the function call and not the delay between the time the message was sent and the time the message was received. By doing so we are able to eliminate the effects of having inconsisent memory available (that will be tested in the following experiment). To achieve this, two user processes were defined, the first sends a message of a fixed size then releases the processor, the second would then receive that message and again release the processor. This process would then repeat a total of 498 times. Data was collected across 5 machines, with 498 calls to each function.

In this single send, single receive condition, a total of 2490 trials were taken, and the data showed no variation for either sending messages or receiving messages. The mean run time of sending a message is 399ns. Due to

33
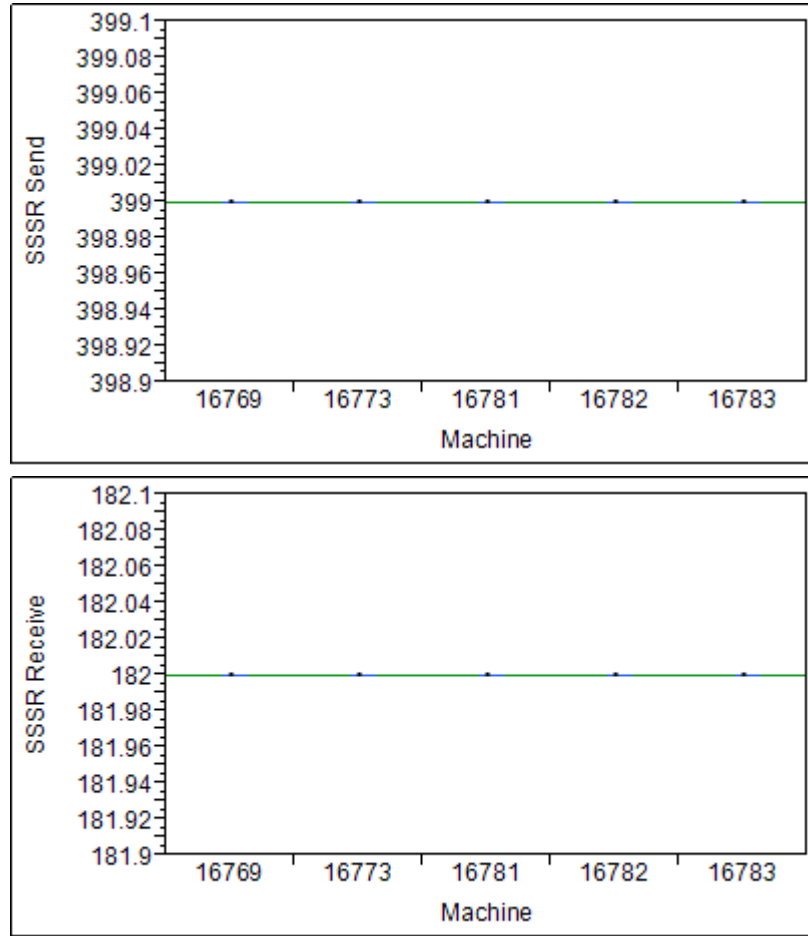
Figure 8.4: Top: Sending messages; Bottom: Receiving messages

the lack of variation on or between each machine this is also the value of the lower and upper bound of the 95% confidence interval. The same applies for receiving the message, which is twice as fast at 182ns.

8.2.2 Cumulative Message Sending

The cumulative message sending condition, as the name implies, had messages accumulate before any were received. The purpose of this is to the see the relation between the number of messages sent (and not yet received) with the time it takes to send a new message. In order to test this, 50 messages

were sent from within a single user processes. Once finished, another process would take over and receive those 50 messages. The results from 5 different machines are displayed in the Figure.
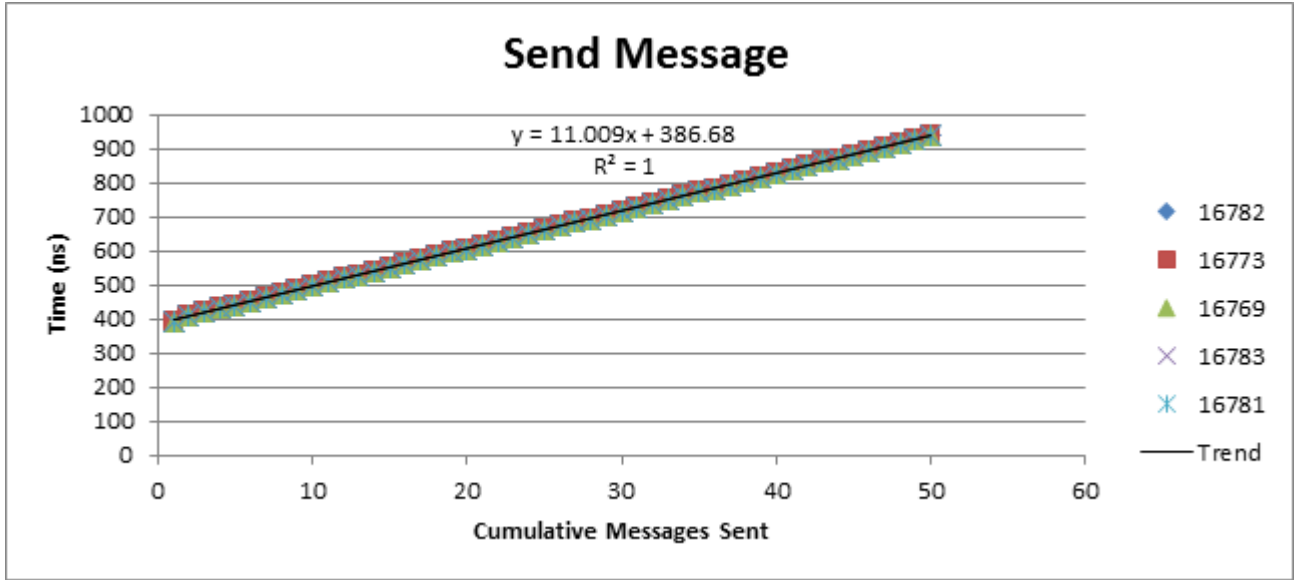
Figure 8.5: Cumulative message sending

Similar to the single send, single receive condition in 8.2.1, the results were the same across all machines. A linear trend line is superimposed onto the results. The trend displayed has an R2 value of 1 showing that this linear trend is very significant. This correlates with the expected O(n) complexity of the send method described in 2.4.1. The trend shows an increase of 11ns per additional message sent.

The message receiving process of this experiment was also measured. For all 250 messages received, the time taken was consistent at 182ns. This is as well expected due to the O(1) complexity of the receive message method describe in 2.4.2.

8.2.3   Effects of Message Length

Another factor that was tested is the length of the message being sent. Depending on the implementation for the kernel it is possible that in some cases a longer message will take longer to send or receive. To test this, messages

35

were continuously sent and received in the same manner as the single send, single receive condition in 8.2.1 with the addition of an accumulating message (ie the first message sent was a, the second message aa and so forth). This was tested up until a message of 50 characters was sent. As shown in the two figures, the message length had no effect on the time taken to send or receive messages.
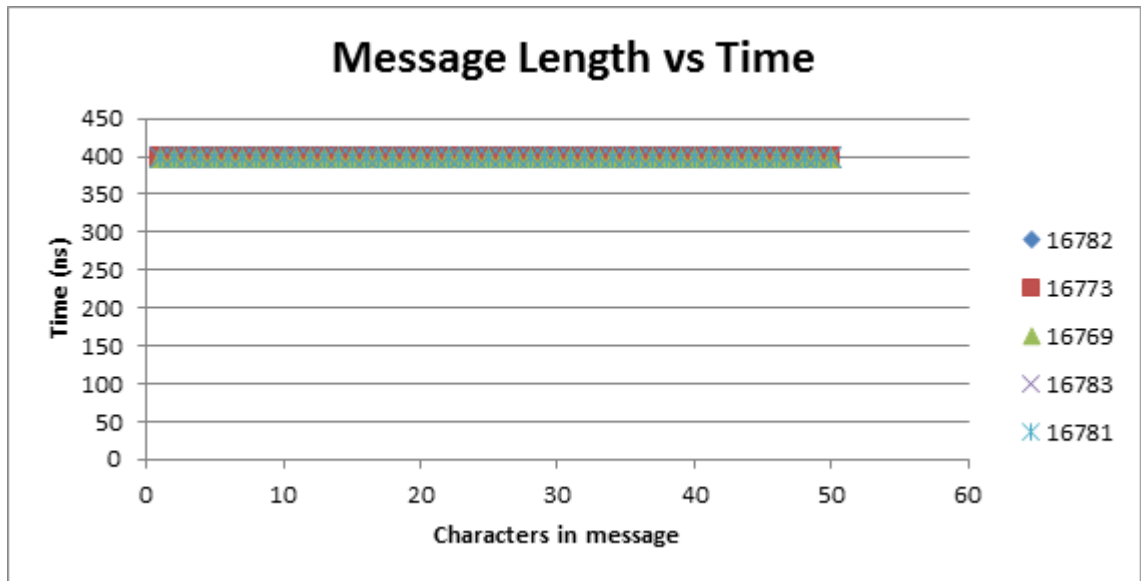
Figure 8.6: Sending messages with variable message length

Appendix A

Raw Measurement Data

A.1   Code Used To Collect Timing Data

```c
static void proc1(void)
{
    timer_test_init();
    static const int memory_block_count = 50;
    static void* memory_blocks[memory_block_count];

    int start_time = 0;
    int end_time = 0;
    LOG("Started single request single release");
    for (int i = 0; i < 1000; i++) {
        start_time = s_test_timer->TC;
        memory_blocks[0] = request_memory_block();
        end_time = s_test_timer->TC;
        LOG("Request: %d", end_time - start_time);
        start_time = s_test_timer->TC;
        release_memory_block(memory_blocks[0]);
        end_time = s_test_timer->TC;
        LOG("Release: %d", end_time - start_time);
    }
    LOG("Finished single request single release");
    LOG("Started requesting");
    for (int i = 0; i < memory_block_count; i++) {
        start_time = s_test_timer->TC;
        memory_blocks[i] = request_memory_block();
        end_time = s_test_timer->TC;
        LOG("%d", end_time - start_time);
    }
```

```
28   LOG("Finished requesting");
     test_results[REQUEST_MEMORY_TEST] = 1;
30   LOG("Started releasing");
     for (int i = 0; i < memory_block_count; i++) {
32     start_time = s_test_timer->TC;
       release_memory_block(memory_blocks[i]);
34     end_time = s_test_timer->TC;
       LOG("%d", end_time - start_time);
36   }
     LOG("Finished releasing");
38
     test_results[RELEASE_MEMORY_TEST] = 1;
40   set_process_priority(1, USER_PROCESS_PRIORITY_LOWEST);
     set_process_priority(2, USER_PROCESS_PRIORITY_HIGH);
42   while (1) {
       s_iteration_count++;
44     release_processor();
     }
46 }

48 // ...

50 static int const required_messages = 50;

52 static void proc3(void)
   {
54   LOG("Started receiving messages");
     for (int i = 0; i < required_messages; i++) {
56     int start_time = 0;
       int end_time = 0;
58     start_time = s_test_timer->TC;
       void* message = receive_message(NULL);
60     end_time = s_test_timer->TC;
       LOG("Receive: %d", end_time - start_time);
62     release_memory_block(message);
     }
64   LOG("Finished receiving messages");
     set_process_priority(5, USER_PROCESS_PRIORITY_HIGH);
66   set_process_priority(4, USER_PROCESS_PRIORITY_LOWEST);
     set_process_priority(3, USER_PROCESS_PRIORITY_LOWEST);
68   release_processor();
   }

70
   static void proc4(void)
72 {
```

```c
     LOG("Started sending messages");
74   for (int i = 0; i < required_messages; i++) {
       struct msgbuf* message_envelope = (struct msgbuf*)
     request_memory_block();
76     message_envelope->mtype = 10;
       strcpy(message_envelope->mtext, test_phrase);
78     int start_time = 0;
       int end_time = 0;
80     start_time = s_test_timer->TC;
       send_message(3, message_envelope);
82     end_time = s_test_timer->TC;
       LOG("Send: %d", end_time - start_time);
84   }
     LOG("Finished sending messages");
86   set_process_priority(3, USER_PROCESS_PRIORITY_HIGH);
     set_process_priority(4, USER_PROCESS_PRIORITY_LOWEST);
88   release_processor();
   }

90
   int sample_size = 500;
92 static void proc5(void)
   {
94   set_process_priority(6, USER_PROCESS_PRIORITY_HIGH);
     void* message = receive_message(NULL);
96   release_memory_block(message);

98   while (true) {
       int start_time = 0;
100    int end_time = 0;
       start_time = s_test_timer->TC;
102    void* message = receive_message(NULL);
       end_time = s_test_timer->TC;
104    LOG("Receive: %d", end_time - start_time);
       release_memory_block(message);
106    release_processor();
     }
108 }

110 static void proc6(void)
   {
112  LOG("Started send-receive together");
     for (int i = 0; i < sample_size; i++) {
114    struct msgbuf* message_envelope = (struct msgbuf*)
     request_memory_block();
       message_envelope->mtype = 10;
```

```
116        strcpy ( message_envelope−>mtext , test_phrase ) ;
           int start_time = 0;
118        int end_time = 0;
           start_time = s_test_timer −>TC;
120        send_message ( 5 , message_envelope ) ;
           end_time = s_test_timer −>TC;
122      LOG( "Send : %d" , end_time − start_time ) ;
           release_processor ( ) ;
124    }
     LOG( "Finished send−receive together" ) ;
126
     LOG( "Started variable message length together" ) ;
128    for (int i = 0; i < sample_size ; i++) {
           struct msgbuf∗ message_envelope = (struct msgbuf∗)
         request_memory_block ( ) ;
130      message_envelope−>mtype = 10;
           for (int j = 0; j < i ; j++) {
132          message_envelope−>mtext[ j ] = ∗"a" ;
           }
134        int start_time = 0;
           int end_time = 0;
136        start_time = s_test_timer −>TC;
           send_message ( 5 , message_envelope ) ;
138        end_time = s_test_timer −>TC;
         LOG( "Send : %d" , end_time − start_time ) ;
140        release_processor ( ) ;
       }
142    LOG( "Finished variable message length together" ) ;
     }
```

A.2 Data Analysis For Single Request Single Release

A.2.1 Request

Summary of Fit

Rsquare	2.63E-05
Adj Rsquare	-0.00077
Root Mean Square Error	11.43196
Mean of Response	1229.237
Observations (or Sum Wgts)	5000

Analysis of Variance

Source	DF	Sum of Square	Mean Square	F Ratio	Prob ¿ F
Machine	4	17.14	4.285	0.0328	0.9979
Error	4995	652795.5	130.69	-	-
C. Total	4999	652812.7	-	-	-

Means of Oneway Anova

Machine	Number	Mean	Std Error	Lower 95%	Upper 95%
16769	1000	1229.29	0.36151	1228.6	1230
16773	1000	1229.29	0.36151	1228.6	1230
16781	1000	1229.17	0.36151	1228.5	1229.9
16782	1000	1229.17	0.36151	1228.5	1229.9
16783	1000	1229.27	0.36151	1228.6	1230

Means and Std Deviations

Machine	Number	Mean	Std Dev	Std Err Mean	Lower 95%	Upper 95%
16769	1000	1229.29	11.7342	0.37107	1228.6	1230
16773	1000	1229.29	11.7342	0.37107	1228.6	1230
16781	1000	1229.17	11.0616	0.3498	1228.5	1229.9
16782	1000	1229.17	11.0616	0.3498	1228.5	1229.9
16783	1000	1229.27	11.5476	0.36517	1228.6	1230

A.2.2 Release

Summary of Fit

Rsquare	3.79E-05
Adj Rsquare	-0.00076
Root Mean Square Error	6.8087
Mean of Response	405.4416
Observations (or Sum Wgts)	5000

Analysis of Variance

Source	DF	Sum of Squares	Mean Square	F Ratio	Prob ¿ F
Machine	4	8.79	2.1968	0.0474	0.9958
Error	4995	231560.2	46.3584	-	-
C. Total	4999	231569	-	-	-

Means for Oneway Anova

Machine	Number of Trials	Mean	Std Error	Lower 95%	Upper 95%
16769	1000	405.416	0.21531	404.99	405.84
16773	1000	405.416	0.21531	404.99	405.84
16781	1000	405.436	0.21531	405.01	405.86
16782	1000	405.416	0.21531	404.99	405.84
16783	1000	405.524	0.21531	405.10	405.95

Means And Std Deviations

Machine	Number	Mean	Std Dev	Std Err Mean	Lower 95%	Upper 95%
16769	1000	405.416	6.56765	0.20769	405.01	405.82
16773	1000	405.416	6.56765	0.20769	405.01	405.82
16781	1000	405.436	6.90519	0.21836	405.01	405.86
16782	1000	405.416	6.56765	0.20769	405.01	405.82
16783	1000	405.524	7.39649	0.23390	405.07	405.98

A.3 Data Analysis for Single Send Single Receive

A.3.1 Send

Summary of Fit

Rsquare	-
Adj Rsquare	-
Root Mean Square Error	0
Mean of Response	399
Observations (or Sum Wgts)	2490

Analysis of Variance

Source	DF	Sum of Square	Mean Square	F Ratio	Prob ¿ F
Machine	4	0	0	-	-
Error	2485	0	0	-	-
C. Total	2489	0	-	-	-

Means of Oneway Anova

Machine	Number	Mean	Std Error	Lower 95%	Upper 95%
16769	498	399	0	399	399
16773	498	399	0	399	399
16781	498	399	0	399	399
16782	498	399	0	399	399
16783	498	399	0	399	399

Means and Std Deviations

Machine	Number	Mean	Std Dev	Std Err Mean	Lower 95%	Upper 95%
16769	498	399	0	0	399	399
16773	498	399	0	0	399	399
16781	498	399	0	0	399	399
16782	498	399	0	0	399	399
16783	498	399	0	0	399	399

A.3.2 Receive

Summary of Fit

Rsquare	-
Adj Rsquare	-
Root Mean Square Error	0
Mean of Response	182
Observations (or Sum Wgts)	2490

Analysis of Variance

Source	DF	Sum of Square	Mean Square	F Ratio	Prob ¿ F
Machine	4	0	0	-	-
Error	2485	0	0	-	-
C. Total	2489	0	-	-	-

Means of Oneway Anova

Machine	Number	Mean	Std Error	Lower 95%	Upper 95%
16769	498	182	0	182	182
16773	498	182	0	182	182
16781	498	182	0	182	182
16782	498	182	0	182	182
16783	498	182	0	182	182

Means and Std Deviations

Machine	Number	Mean	Std Dev	Std Err Mean	Lower 95%	Upper 95%
16769	498	182	0	0	182	182
16773	498	182	0	0	182	182
16781	498	182	0	0	182	182
16782	498	182	0	0	182	182
16783	498	182	0	0	182	182