

RTX Project Report

Xiang, Dian

20431601

`dxiang@uwaterloo.ca`

Justin McGirr

20413625

`jmcgirr@uwaterloo.ca`

Adrian Cheung

20421743

`a32cheun@uwaterloo.ca`

Aaron Morais

20413440

`aemorais@uwaterloo.ca`

April 6, 2014

Abstract

Contents

I	Introduction	6
II	Awesome RTX	7
1	Global Variables	8
1.1	Description	8
1.2	Heap Data Structures	8
1.3	Process Scheduler Data Structures	9
1.3.1	Ready and Blocked Priority Queues	10
1.3.2	Priorities	11
1.3.3	Message Passing	12
2	Kernel API	13
2.1	Description	13
2.2	Memory Management	13
2.2.1	Description	13
2.2.2	Requesting Memory	14
2.2.3	Releasing Memory	15
2.3	Processor Management	16
2.4	Interprocess Communication	16
2.4.1	Send Messages	17
2.4.2	Receive Messages	18
2.4.3	Receive Messages	19
2.5	Process Priority	19
2.5.1	Set Process Priority	19
2.5.2	Get Process Priority	20

3	Interrupts and Their Handler and Processes	21
3.1	Description	21
3.2	UART	21
3.3	TIMER	21
4	System and User Processes	22
4.1	System Processes	22
4.1.1	Description	22
4.1.2	Null Process	22
4.1.3	CRT Display Process	23
4.1.4	Keyboard Command Decoder Process	23
4.2	User Processes	25
4.2.1	Wall Clock Process	25
4.2.2	‘funProcess’	25
4.2.3	‘schizophrenicProcess’	25
4.2.4	‘fibProcess’	25
4.2.5	‘memoryMuncherProcess’	25
4.2.6	‘releaseProcess’	25
5	Initialization	26
5.1	Description	26
5.2	Memory Layout	26
5.3	Process Mailbox	26
6	Major Design Changes	27
6.1	Description	27
6.2	Heap	27
6.3	Scheduler	27
6.4	What We Learned	27
III	Testing and Analysis	28
7	Testing	29
7.1	Description	29
7.2	Theoretical Analysis	29
7.3	Measurements	29

8	Timing Analysis	30
8.1	Requesting and Releasing Memory	30
8.1.1	Single Request Single Release	30
8.1.2	Cumulative Memory Requests	31
8.2	Sending and Receiving Messages	32
8.2.1	Single Send, Single Receive	32
8.2.2	Cumulative Message Sending	32
8.2.3	Affects of Message Length	33
A	Raw Measurement Data	34
A.1	Trial Information	34
A.2	Function Runtime Profiling	35

List of Algorithms

List of Figures

Part I

Introduction

Part II

Awesome RTX

Chapter 1

Global Variables

1.1 Description

There are various amounts of variables used throughout the RTX. This section describes the global variables that are used to accomplish tasks in the RTX. Other global variables such as the global process IDs can be seen in the appendix TODO. There are three major sections that required global variables and data structures: memory management in the heap, scheduler, and user processes.

1.2 Heap Data Structures

The RTX provides functionality to request and release memory from the heap, which is shared and stored in the RAM of the board. There is one main data structure that stores each memory block.

```
1 typedef struct HeapBlockHeader {  
    int source_pid;  
3    int dest_pid;  
    unsigned int send_time;  
5    struct HeapBlock* p_next;  
    struct HeapBlock* p_next_usr;  
7 } HeapBlockHeader;  
  
9 #define HEAP_BLOCK_SIZE 128  
  
11 typedef struct HeapBlock {  
    HeapBlockHeader header;
```

```

13  byte          data[HEAP_BLOCK_SIZE];
    } HeapBlock;

```

The HeapBlock structure stores a header and the content of the block. When a memory is requested, users are given the data with a adjustable size of HEAP_BLOCK_SIZE and does not have knowledge of the header. Helper functions in the RTX turn the user block back into kernel block by adjusting the pointer of the block. Each header contains information about message passing (if the block is used for this purpose) and a pointer to the next block (for kernel memory management purposes). This global data structure is used in the following sections:

1. **Memory Management** - this section uses this data structure to give and receive back heap memory blocks
2. **Process Message Passing** - this section uses this data structure pass messages between processes. Message envelope are implemented as HeapBlocks.
3. **Timer I-Process** - The process uses message passing to send delayed messages.
4. **UART I-Process** - The process uses message passing to process input and output characters.
5. **CRT I-Process** - The process receives messages and passes message for terminal output.
6. **KCD Process** - The KCD uses message passing to CRT for display and user processes for processing.

1.3 Process Scheduler Data Structures

The RTX has a fixed-priority based scheduler that acts as a uniprocessor system. Context switching is required and the following data structures are used for this purpose.

1.3.1 Ready and Blocked Priority Queues

On initialization, each in memory process is given a process control block (PCB). The PCB contains data on the process such as its process ID, stack pointer, process state, and priority that the kernel will use for scheduling. The blocked and ready priority queues of PCBs keep track of which processes are blocked and ready for execution. Each process is in one of the 5 states at all time listed in `ProcessState`. More details about the states can be found in section TODO. Processes that are in state `PROCESS_STATE_READY` are on the ready queue. Processes in state `PROCESS_STATE_BLOCKED` are in the blocked queue. There is also an implicit queue for processes blocked on message but we won't go into that in this section.

```
1 typedef enum {
2     PROCESS_STATE_NEW           = 0,
3     PROCESS_STATE_READY        = 1,
4     PROCESS_STATE_RUNNING      = 2,
5     PROCESS_STATE_BLOCKED      = 3,
6     PROCESS_STATE_BLOCKED_ON_MESSAGE = 4,
7 } ProcessState;

9 typedef struct PCB {
10     // Stack pointer
11     U32*      sp;

12     ProcessID  pid;
13     ProcessState state;
14     ProcessPriority priority;

15     struct PCB* p_next;
16     // Incoming messages, waiting to be processed.
17     HeapBlock* message_queue;
18 } PCB;

21 PCB* g_ready_process_priority_queue[PROCESS_PRIORITY_NUM] = {
22     NULL};
23 PCB* g_blocked_process_priority_queue[PROCESS_PRIORITY_NUM];
```

PCBs and the priority queues are used in the following sections:

1. **Memory Management** - When a process requests memory without available memory blocks, the memory management unit must add the process to a blocked queue. When a memory block is released, the pro-

cess with the highest priority in the blocked queue (if any) is moved into the ready queue. Thus, the memory management unit must have access to the PCBs, and blocked and ready queues in order to accomplish these task.

2. **Process Management** - The process management unit is the one that schedules processes and keep track of the process priorities. Ready processes are taken from the ready queue to be ran if the current process is blocked or finished its execution. The process management unit also takes care of updating the two priority queues when the priority of a process has been changed.
3. **UART Process** - The UART Process needs access to the PCB. The PCBs also contain a mailbox for messages. The UART displays any incoming messages to the terminal display. Thus, it needs to know the PCB structure in order to gain access to the mailbox.

1.3.2 Priorities

The RTX is based on a fixed-priority scheduler with 5 user process priorities and 2 system process priorities given below. Users are given priorities `PROCESS_PRIORITY_HIGH` to `PROCESS_PRIORITY_LOWEST`. `PROCESS_PRIORITY_NULL_PROCESS` is given to the null process and `PROCESS_PRIORITY_SYSTEM_PROCESS` are given to critical processes such as the KCD, CRT, timer I-process, and UART I-process.

```

1 typedef enum {
2     PROCESS_PRIORITY_INVALID          = 0,
3     PROCESS_PRIORITY_SYSTEM_PROCESS = 1,
4     PROCESS_PRIORITY_HIGH             = 2,
5     PROCESS_PRIORITY_MEDIUM          = 3,
6     PROCESS_PRIORITY_LOW             = 4,
7     PROCESS_PRIORITY_LOWEST          = 5,
8     PROCESS_PRIORITY_NULL_PROCESS    = 6,
9     PROCESS_PRIORITY_UNSCHEDULABLE  = 7,
10
11     PROCESS_PRIORITY_NUM              = 8
12 } ProcessPriority;
13
14 typedef enum {
15     USER_PROCESS_PRIORITY_HIGH        = 0,
16     USER_PROCESS_PRIORITY_MEDIUM     = 1,

```

```

17 USER_PROCESS_PRIORITY_LOW           = 2,
   USER_PROCESS_PRIORITY_LOWEST      = 3,
19
   USER_PROCESS_PRIORITY_NUM         = 4,
21 } UserProcessPriority;

```

This priority structure is used by the process management unit to schedule and block processes based on their priorities. The RTX provides users with `UserProcessPriority` while keeping an internal structure of `ProcessPriority`.

1.3.3 Message Passing

Message passing is a way for interprocess communication provided by the RTX. Messages are created in the form of a `msgbuf` structure, which includes the type and content. Some processes such as the CRT and KCD will require a certain type of message to be sent before the message can be processed correctly.

```

1 typedef enum {
   MESSAGE_TYPE_KCD_KEYPRESS_EVENT      = 0,
3   MESSAGE_TYPE_KCD_COMMAND_REGISTRATION = 1,
   MESSAGE_TYPE_CRT_DISPLAY_REQUEST     = 2,
5   MESSAGE_TYPE_WALL_CLOCK              = 3,
   MESSAGE_TYPE_COUNT_REPORT             = 4,
7   MESSAGE_TYPE_WAKEUP_10               = 5,

9   MESSAGE_TYPE_NUM                     = 6,
} MessageType;
11
12 struct msgbuf {
13   MessageType mtype;
   char mtext[HEAP_BLOCK_SIZE - sizeof(MessageType)];
15 };

```

The message passing data structure is evident in the KCD, wall clock process, and all I-processes. It can also be used in all user processes.

Chapter 2

Kernel API

2.1 Description

This section describes the kernel API that is available to users in the RTX. It will only go into details of how to use each API function call and states of different scenarios. Details of the implementation can be found in section TODO or from the raw code.

2.2 Memory Management

2.2.1 Description

The RTX provides the utility of simple memory management of the heap. The main memory on the board (TODO: name the board) is divided into sections of the RTX Image, PCB data, the heap, and process stacks seen in figure blah. *** TODO insert figure blah here ***

Board LPC17xx (TODO insert the real name) does not have the hardware to support virtual memory. Thus, a fixed memory management scheme is used with variable block sizes. The OS kernel is always loaded into main memory. When the OS boots, a user stated number of PCBs and process stack are allocated. The OS also allocates memory for system processes such as the Keyboard Command Decoder Process (KCD) and the CRT Display Process. After allocation of PCBs and process stacks, the remaining memory is used for the heap, which is shared between all processes. This section will be focused on the memory management of the heap.

The heap is further divided into a variable number of blocks. Each block contains a header (HeapBlockHeader) and 128 bytes of data. Depending on the number of user and system processes, the number of available heap memory will vary. The size of the data can vary from 128 bytes but must be set at compile time. The OS supports two kernel calls that gives access to these memory blocks.

2.2.2 Requesting Memory

```
void *request_memory_block();
```

The first functionality supported by the OS is the ability to request memory. This call gives back a pointer to a memory block in the heap. The size of the memory block is defined by the constant `HEAP_BLOCK_SIZE`, which defaults to 128 bytes. These blocks are used for storing local variables or as envelopes for interprocess communication (described in section TODO). The user must cast the memory block to the proper type for his or her own use. An example of the usage is provided below, which stores the numbers 0-31 in an array of size 32.

```
1 void user_process() {  
    int size = 32;  
3  
    int* array = (int*)request_memory_block();  
5  
    for( int i = 0; i < size; i++ ) {  
7        array[i] = i;  
    }  
9  
    // ...  
11  
    release_memory_block( (void*)array );  
13 }
```

All processes will share the same heap memory pool. Thus, this primitive will block the process if the OS does not have anymore memory blocks to give out at the time, thus effecting the execution of the process. In that case, it will only be unblocked there is a new memory block available and if it has the highest priority on the list of processes waiting for a memory block. When using the memory block, the user must be aware of writing past the

heap block size. The OS does not check for segmentation faulting. Thus, undefined behavior may occur. Also, the user must remember to release the memory block or memory leakage will occur.

Requesting memory is reliant on calls to the heap to find an available memory block to allocate. The heap keeps a list of the status of each memory block to determine whether or not a block of memory can be requested. The following is the implementation used by the heap:

```
function FINDFREEBLOCK
  for  $i = 0$  to TotalNumberOfBlocks do
    if FreeSpaceBitmap[i] is a free block then
      return i
    end if
  end for
  return not found
end function
```

The lookup time of the FreeSpaceBitmap is $O(1)$ as it is a simple array lookup and a boolean evaluation. Finding a free block has its dependencies on the total number of blocks. The best case for requesting memory is $O(1)$ when the first block in the heap is free. Since this is a linear search, on average the complexity is $O(n)$ where n is the total number of blocks.

2.2.3 Releasing Memory

```
1 int release_memory_block(void *memory_block);
```

The second functionality supported by the OS is the ability to release memory. This is a non-blocking call that returns the memory block back to the OS. This should be called when a message is received and not passed on or when the process is done using the memory block. If any process is blocked on memory, it will be unblocked and put to the ready queue. If the current process has lower priority than the unblocked process, then the current process will be preempted and the higher priority process will be executed instead. An example of this can be seen in figure CODE1 TODO on line 13.

In order to release the memory, `release_memory_block` will need to rely on the implementation of the heap. As stated in 2.2.2, the heap keeps a list

of all available memory blocks. In order to release a memory block, the heap will set that block in the list as a free block. Releasing the block thus has $O(1)$ complexity as it only relies on an array lookup. The implementation is as follows:

```
function RELEASEMEMORY(MemoryBlock)
    FreeSpaceBitmap[position of MemoryBlock]  $\leftarrow$  true
end function
```

2.3 Processor Management

```
1 int release_processor();
```

The OS manages processes as though it is on a uniprocessor. A priority based scheme with context switching is used for scheduling processes. A process can voluntarily release the processor to the OS at any time during its execution. If there are errors in the call, it will return an error without releasing the processor to the OS. If there are no errors and the invoking process is ready to execute, it is put to the end of the ready queue of its priority. If there are no other processes of equal or higher priority, the process will be chosen to execute again. However, if there is, another process may be chosen for execution. Below is an example which prints an increasing number and releases the processor at every turn.

```
1 void user_process() {
    static num = 0;
3   while(1) {
        print(num++);
5       release_processor();
    }
7 }
```

2.4 Interprocess Communication

Apart from heap memory, processes do not share information. Thus, communication between processes is done through message-based interprocess communication (IPC). Details of the internal layout of process mailboxes

can be found in section TODO. The RTX gives three primitives to carry out this task, one for sending, one for delayed sending, and one for receiving.

2.4.1 Send Messages

```
1 int send_message(int process_id, void *message_envelope);
```

A process can compose a message envelope to be sent to another process. Memory for envelope message must be requested from the RTX using the request_memory() routine. The envelope consists of a type (mtype) and the message data (mtext) which must be filled in as seen below. The predefined message types are used for the KCD, CRT, and Wall Clock process, which are built into the RTX (see section TODO for more information). The message data has a predefined size which is a MessageType smaller than the HEAP_BLOCK_SIZE. Thus, any message that are longer will exhibit undefined behavior. The process ID of the receiving process must also be known ahead of time in order to use send_message.

```
1 typedef enum {
    MESSAGE.TYPE.KCD.KEYPRESS_EVENT      = 0,
3    MESSAGE.TYPE.KCD.COMMAND.REGISTRATION = 1,
    MESSAGE.TYPE.CRT.DISPLAY.REQUEST     = 2,
5    MESSAGE.TYPE.WALL.CLOCK              = 3,
    MESSAGE.TYPE.USER.DEFINED             = 4,
7
    MESSAGE.TYPE.NUM                      = 5,
9 } MessageType;

11 struct msgbuf {
    MessageType mtype;
13    char mtext[HEAP_BLOCK_SIZE - sizeof(MessageType)];
};
```

The primitive returns a status which validates the, message, receiver process ID, and ready queue. User processes is allowed to send a message to any user processes or system processes such as the KCD. If the receiving process is currently blocked on receiving message, this call will add the message to the process' message box and put it back on the message queue. The current process continues to execute unless the receiving process has a higher

priority. In that case, the current process will be preempted and put to the back of the ready queue. Thus, this primitive may effect the execution of the process.

In order to keep track of messages that have been sent, the messages are stored in a queue, specifically a linked list. Messages are added to the end of the linked list resulting in an $O(n)$ insertion complexity. It is possible to have this run in $O(1)$ time by having the inserting take place at the front of the linked list, however this is a trade-off for receiving messages. In the current implementation, received messages will be taken from the front of the linked list having $O(1)$ complexity. If the change was to be made to optimize sending messages, then this would slow down receiving messages to $O(n)$.

2.4.2 Receive Messages

```
void *receive_message(int *sender_id);
```

A process can use this primitive to receive messages from other processes. Unless the sender_id is NULL, the sender of the message will be written to the sender_id parameter. The current process will check its mailbox for any incoming messages. If there are no messages in its mailbox, the routine will block the process and select another process for execution. Execution of the process will occur again if a message is sent from another process and this process has the highest priority in the ready queue. The message_envelope are heap memory underneath. Thus, it is required that unless the process passes the message onto another process, the final receiving process must release the message_envelope block after its usage. An example of send_message and receive_message is shown below.

```
1 void process_send() {  
    while(1) {  
3         // Request one block for process_receive to use  
        struct msgbuf* message_envelope  
5         = (struct msgbuf*)request_memory_block();  
        message_envelope->mtype = MESSAGE_TYPE_USER_DEFINED_1;  
7        strcpy(message_envelope->mtext, "process message");  
        send_message(PROCESS_RECEIVE_ID, message_envelope);  
9  
        // Send another block for process_receive to send to CRT  
11       struct msgbuf* message_envelope  
        = (struct msgbuf*)request_memory_block();
```

```

13     message_envelope->mtype = MESSAGE_TYPE_USER_DEFINED_2;
14     strcpy(message_envelope->mtext, "print this");
15     send_message(PROCESS_RECEIVE_ID, message_envelope);
16 }
17 }
18
19 void process_receive() {
20     while(1) {
21         struct msgbuf* message
22             = (struct msgbuf*)receive_message(NULL);
23         if( message->mtype = MESSAGE_TYPE_USER_DEFINED_2 ) {
24             // Send to CRT for printing; do not release memory
25             message->mtype = MESSAGE_TYPE_CRT_DISPLAY_REQUEST;
26             send_message(PROCESS_ID_CRT, message);
27         } else {
28             // Do something with the message then release it
29             release_memory_block( (void*)message );
30         }
31     }
32 }

```

2.4.3 Delayed Messages

```

int delayed_send(int process_id, void *message_envelope, int
delay);

```

This primitive is very similar to the primitive of `send_message` in section 2.4.1 with the addition of a delay. Instead of sending the message immediately, the message will be sent delay number of milliseconds. The `message_envelope` is constructed in the same way. The `process_id` will be the receiving process. The execution of the process may be preempted after a the delay period if the receiving process was blocked on `receive_message` and has a higher priority.

2.5 Process Priority

The RTX schedules processes based on a fixed priority scheme. Thus, the RTX provides the utility to change priorities and to get the process priority of any user processes.

2.5.1 Set Process Priority

```
1 int set_process_priority(int process_id, int priority);
```

This primitive allows user processes to set the priorities of other user processes. User processes are not allowed to set the priorities of any system processes. System processes, however, are allowed to set the priorities of other system or i-processes. The primitive validates if the current process is allowed to set the priority of the process with process_id and whether it's a valid process_id. It also checks the validity of the priority. An error status of RTX_ERR is given back if the validation does not pass. If properly set, a status of RTX_OK is returned. The caller of this primitive is not blocked but can be preempted if the priority of the set process is higher. Otherwise, the process continues to execute. An example usage is shown in section [2.5.2](#)

2.5.2 Get Process Priority

```
1 int get_process_priority(int process_id);
```

Given a process_id, this primitive gives back the priority of any processes including any system or i-processes to a user process. A invalid process_id will result in a RTX_ERR status. Otherwise, a RTX_OK is returned to the caller. An example usage of both get and set process priority is shown below.

```
1 // Assume this process has a process ID of 3
2 // and there is another user process with ID of 1
3 void user_process() {
4     int process_priority_3 = get_process_priority(3);
5     if( process_priority_3 == USER_PROCESS_PRIORITY_MEDIUM ) {
6         // Make process 1 have a higher priority
7         // This call will preempt
8         int status = set_process_priority(
9             1, USER_PROCESS_PRIORITY_HIGH );
10        if( status == RTX_ERR ) {
11            release_processor();
12        }
13    }
14 }
```

Chapter 3

Interrupts and Their Handler and Processes

3.1 Description

3.2 UART

3.3 TIMER

Chapter 4

System and User Processes

4.1 System Processes

4.1.1 Description

The RTX consists of system processes that sit on top of the kernel. These system processes usually have higher priority than user processes because they carry out tasks for the OS that the kernel cannot do otherwise. System processes in this RTX includes the null process, and I/O processes.

4.1.2 Null Process

The design of the RTX requires that a process should always be executing. Thus, if there are no user or other system processes, the null process executes. The only job of the null process is to keep the processor running. It has the lowest priority of any process using the RTX. If any process is added into the ready queue during the execution of this process, a preemption will occur and the new process will take control of the processor. The final implementation follows the following form.

```
function NULL_PROCESS
    while true do
        end while
end function
```


4.1.3 CRT Display Process

The RTX supports output to terminal display. CRT is a system process that the RTX uses to accomplish the task of display. The CRT process depends on the UART I-Process (details of the UART can be found in section TODO). The CRT is implemented using the following algorithm.

```
function CRT_PROCESS
  while true do
    message = RECEIVE_MESSAGE(NULL)
    if message is not a CRT display request then
      continue
    end if
    SEND_MESSAGE(PROCESS_ID_UART, message)
    Set interrupt bits
  end while
end function
```

The CRT display process accepts message requests from any process of type MESSAGE_TYPE_CRT_DISPLAY_REQUEST. This tells the CRT to send the message for display to the UART I-process, which interacts with the hardware registers to output to terminal. Since the UART process is an interrupt process, the interrupt bit must be set in order for the process to run. Thus, after sending the message to UART's mailbox, the CRT process enables the interrupt bits for the I-process to trigger. The UART I-process passes the message character by character to the hardware to display to the terminal.

4.1.4 Keyboard Command Decoder Process

The RTX provides functionality to forward user input commands to a particular process. The keyboard command decoder process (KCD) handles two tasks: giving input to the CRT for terminal display and forwarding certain inputs to different processes for further processing. For the first task, all keyboard inputs are filtered by the KCD. The UART I-process is called when a key is first pressed. The UART I-process (explained further in section TODO) then passes a message to the KCD telling of type key press. In the first case, the KCD will take this message and pass it on to the CRT for terminal display.

In the latter case, processes can register a command with the KCD. Each

command registration consists of a single capital letter of the alphabet. If a command is taken, the KCD will not register the command and exit with an error. After a valid command has been registered, user input with the pattern "%" followed by a upper case character will be parsed by the KCD. If user types in "%" followed by a registered character, the KCD will buffer the input until the user presses the return key. In that scenario, the KCD will pass it onto the registered process. The following algorithm of the KCD explains this in more detail.

```

function PROCESS_INPUT(message)
    if There a "%" character in the buffer then
        if message.Content is a command character then
            add to buffer
        else
            clear buffer
        end if
    else if message.Content is "%" then
        add to buffer
    else if message.Content is enter and buffer is not empty then
        send buffer to registered command
    end if
    Send message to CRT to display
end function

function PROCESS_REGISTER(message)
    if message.Command has not been taken then
        register message.Command
    else
        error
    end if
end function

function KCD_PROCESS
    initialization
    while true do
        message = RECEIVE_MESSAGE(NULL)
        if message type is key press input then
            PROCESS_INPUT(message)
        else if message type is command registration then

```

```

        PROCESS_REGISTER(message)
    end if
end while
end function

```

An example of this usage is the wall clock process (described in more detail in section 4.2.1). The wall clock process registers the character 'W' as a command. The registration will be done by passing a message to the KCD of command registration type. This message is received by the KCD. If the 'W' command has not been registered, it becomes registered by the wall clock process. Meanwhile, the user types in random characters. Any character that is not '%' is passed by the KCD to the CRT for display. Once a '%' is received, the KCD checks the next character to see if it's a registered command. If its not a registered command, the KCD clears the buffer and starts at the beginning. If it is a registered command, the KCD keeps a buffer of all the input characters. This buffer fills with data and is sent to the wall clock process for processing after the user presses enter. The buffer is only 20 characters long. Thus, any data for a process after 20 characters will not be sent to the registered process on enter.

4.2 User Processes

4.2.1 Wall Clock Process

4.2.2 'funProcess'

4.2.3 'schizophrenicProcess'

4.2.4 'fibProcess'

4.2.5 'memoryMuncherProcess'

4.2.6 'releaseProcess'

Chapter 5

Initialization

5.1 Description

5.2 Memory Layout

5.3 Process Mailbox

Chapter 6

Major Design Changes

6.1 Description

6.2 Heap

6.3 Scheduler

6.4 What We Learned

Part III

Testing and Analysis

Chapter 7

Testing

7.1 Description

7.2 Theoretical Analysis

7.3 Measurements

Chapter 8

Timing Analysis

8.1 Requesting and Releasing Memory

8.1.1 Single Request Single Release

In the first analysis, the `request_memory` and `release_memory` methods were timed. Single request, single release means that a single memory block was requested and the time taken was recorded. Similarly for releasing a memory block. By doing so we are able to keep the heap in the same state, eliminating the effects of cumulative memory requests (which will be tested in section 8.1.2). The method to time this is shown in the following algorithm. Data was collected across 5 machines, with 1000 calls to each function.

```
function SINGLE_REQUEST_SINGLE_RELEASE
  for  $i = 0$  to 1000 do
    RequestStartTime[i]  $\leftarrow$  Curent Time
    memoryBlock  $\leftarrow$  REQUEST_MEMORY_BLOCK
    RequestEndTime[i]  $\leftarrow$  Curent Time
    ReleaseStartTime[i]  $\leftarrow$  Curent Time
    RELEASE_MEMORY_BLOCK(memoryBlock)
    ReleaseEndTime[i]  $\leftarrow$  Curent Time
  end for
end function
```

Based on the analysis, the mean run time for release memory is 405.44ns. The standard error (marked in the green mean diamonds on the figure) is very small, giving a 95% confidence interval ranging from [404.99, 405.84] to

[405.1, 405.95]. The F ratio is incredibly small (0.0474), meaning that there is little to no variance among the different machines and as such it can be concluded that the hardware does not greatly affect the performance of the method.

The results of releasing memory are very similar to that of requesting memory. The mean time taken to release memory is 1229.237ns. The variations among different machines were very small, giving an F Ratio of 0.0328. This shows that the performance is consistent across different hardware. The performance was as well consistent within itself, the fastest machine had a 95% confidence interval of [1228.5, 1229.9], and for the slowest [1228.6, 1230.0]. Releasing memory takes on average a little under three times as long as requesting memory

8.1.2 Cumulative Memory Requests

The cumulative memory requests condition had 50 memory blocks requested and then subsequently released. This differs from experiment in 8.1.1 where the requested memory blocks do not accumulate. The purpose of this experiment is to see the effect of depleting resources on run time. Once again, trials were taken over 5 machines, however, each machine showed the same performance. To achieve this, the following two method calls were made.

```

function CUMULATIVE_REQUEST
  for  $i = 0$  to 50 do
    RequestStartTime[i]  $\leftarrow$  Curent Time
    memoryBlock  $\leftarrow$  REQUEST_MEMORY_BLOCK
    RequestEndTime[i]  $\leftarrow$  Curent Time
  end for
end function

function CUMULATIVE_RELEASE
  for  $i = 0$  to 50 do
    ReleaseStartTime[i]  $\leftarrow$  Curent Time
    memoryBlock  $\leftarrow$  RELEASE_MEMORY_BLOCK
    ReleaseEndTime[i]  $\leftarrow$  Curent Time
  end for
end function

```

Consistent with the Single Request, Single Release experiment, the first memory block requested took 1229ns. A linear trend line is overlaid on top

showing an R^2 value of 1 (a perfect trend). The trend shows that for each additional memory block requested, the run time increase by 17ns. This linear trend is in accordance with the implementation of the request memory method which has $O(n)$ complexity described in [2.2.2](#).

8.2 Sending and Receiving Messages

8.2.1 Single Send, Single Receive

In this experiment, the `send_message` and `receive_message` methods were timed. A single message was sent and the time taken was recorded. Similarly, the time taken to receive that message was timed, specifically the time for the function call and not the delay between the time the message was sent and the time the message was received. By doing so we are able to eliminate the effects of having inconsistent memory available (that will be tested in the following experiment). To achieve this, two user processes were defined, the first sends a message of a fixed size then releases the processor, the second would then receive that message and again release the processor. This process would then repeat a total of 498 times. Data was collected across 5 machines, with 498 calls to each function.

In this single send, single receive condition, a total of 2490 trials were taken, and the data showed no variation for either sending messages or receiving messages. The mean run time of sending a message is 399ns. Due to the lack of variation on or between each machine this is also the value of the lower and upper bound of the 95% confidence interval. The same applies for receiving the message, which is twice as fast at 182ns.

8.2.2 Cumulative Message Sending

The cumulative message sending condition, as the name implies, had messages accumulate before any were received. The purpose of this is to see the relation between the number of messages sent (and not yet received) with the time it takes to send a new message. In order to test this, 50 messages were sent from within a single user processes. Once finished, another process would take over and receive those 50 messages. The results from 5 different machines are displayed in the Figure.

Similar to the single send, single receive condition in [8.2.1](#), the results

were the same across all machines. A linear trend line is superimposed onto the results. The trend displayed has an R^2 value of 1 showing that this linear trend is very significant. This correlates with the expected $O(n)$ complexity of the send method described in [2.4.1](#). The trend shows an increase of 11ns per additional message sent.

The message receiving process of this experiment was also measured. For all 250 messages received, the time taken was consistent at 182ns. This is as well expected due to the $O(1)$ complexity of the receive message method describe in [2.4.2](#).

8.2.3 Affects of Message Length

Another factor that was tested is the length of the message being sent. Depending on the implementation for the kernel it is possible that in some cases a longer message will take longer to send or receive. To test this, messages were continuously sent and received in the same manner as the single send, single receive condition in [8.2.1](#) with the addition of an accumulating message (ie the first message sent was a, the second message aa and so forth). This was tested up until a message of 50 characters was sent. As shown in the two figures, the message length had no effect on the time taken to send or receive messages.

Appendix A

Raw Measurement Data

A.1 Trial Information

Trial	Total Runtime	Notes
1	4.219	Normal (no stress processes)
2	7.754	Wall clock
3	8.487	Normal (no stress processes)
4	6.5	No Memory Muncher or Release Process
5	30.988	Stress processes

A.2 Function Runtime Profiling

Function	Trial	Time (μs)	# of Calls	Average time / call (μs)
k_sendMessage	1	601.58	552	1.090
k_receiveMessage	1	408.22	565	0.723
k_acquireMemoryBlock	1	244.12	294	0.830
k_sendMessage	2	647.44	594	1.090
k_receiveMessage	2	437.78	606	0.722
k_acquireMemoryBlock	2	258.68	320	0.808
k_sendMessage	3	630.99	579	1.090
k_receiveMessage	3	426.83	591	0.722
k_acquireMemoryBlock	3	259.24	321	0.808
k_sendMessage	4	108.80	100	1.088
k_receiveMessage	4	74.44	110	0.677
k_acquireMemoryBlock	4	92.47	123	0.752
k_sendMessage	5	750.63	687	1.093
k_receiveMessage	5	497.09	693	0.717
k_acquireMemoryBlock	5	329.90	447	0.738